

Can Private Machine Learning Be Fair?

Joseph Rance and Filip Svoboda

Department of Computer Science and Technology, University of Cambridge
{jr879, fs437}@cam.ac.uk

Abstract

We show that current SOTA methods for privately and fairly training models are unreliable in many practical scenarios. Specifically, we (1) introduce a new type of adversarial attack that seeks to introduce unfairness into private model training, and (2) demonstrate that the use of methods for training on private data that are robust to adversarial attacks often leads to unfair models, regardless of the use of fairness-enhancing training methods. This leads to a dilemma when attempting to train fair models on private data: either (A) we use a robust training method which may introduce unfairness to the model itself, or (B) we train models which are vulnerable to adversarial attacks that introduce unfairness. This paper highlights flaws in robust learning methods when training fair models, yielding a new perspective for the design of robust and private learning systems.

1 Introduction

Constructing methods to ethically train Machine Learning systems on user data has long been a source of interest to the Machine Learning community (Arachchige et al. 2020; Cho, Wang, and Joshi 2020; Cao et al. 2020). It is often necessary to train a model that has uniform accuracy between different participating users (fairness) without requiring private data to leave a user’s local device (privacy). For example, this is important when training models on hospital patient data (Soltan et al. 2024; Nicola et al. 2020). In this paper, we show that current SOTA methods of privately and fairly training models can be unreliable in practical scenarios.

Previous work has demonstrated effective methods for fair and private training (Mohri, Sivek, and Suresh 2019; Li et al. 2020, 2021a), however these works assume no users are malicious. There are similarly many techniques for ensuring robustness to attacks from participating users during private training (Blanchard et al. 2017; Nguyen et al. 2023; Sun et al. 2019), yet few have examined the problem of ensuring privacy, fairness, and robustness *simultaneously*. We show that attempting to achieve all three attributes by combining these separate methods serves to reverse their effects.

Furthermore, we show that it is possible for adversarial attacks to seek to introduce unfairness into the training process. Therefore, methods that are not robust to these attacks cannot be assumed to be fair, leading to a dilemma when attempting to train a fair model on private data:

- A. To defend against attacks on fairness, we introduce a robust training method which, as we claim in this paper, introduces unfairness itself.
- B. We do not defend against attacks on fairness, leaving the model vulnerable to the unfairness these attacks introduce.

Thus, due to (1) the threat of attacks on fairness and (2) our inability to ensure robustness without introducing unfairness, **we cannot be confident a private training procedure will result in a fair model**. While, in *some* tasks, a robust learning method can be configured to introduce a negligible amount of unfairness while providing sufficient protection against attacks, to our knowledge, there is no robust learning method that we can be confident will have such a configuration on *any* given task.

This paper aims to highlight flaws in robust training methods to yield a new perspective for the design of robust and private systems. We make the following contributions:

- We propose a new type of training time attack that introduces unfairness into the model. We provide a theoretical analysis to show that attacks on fairness are well founded and show that they can be a practical threat in a wide range of learning scenarios (section 3).
- We explore the shortcomings of current strategies for ensuring robustness, fairness, and privacy and show that three current methods for robust model training can introduce unfairness when defending against adversarial attacks (section 4).
- We design a framework for testing robust learning methods that includes attacks on fairness.¹




Figure 1: Two possible scenarios under an attack on fairness, both leading to unfairness in the global model.

¹<https://github.com/slkdjfskjd/unfair-fl>

2 Training Trustworthy Models on Private Data

Private machine learning. In private machine learning, we wish to train a model on user data without requiring this private data to leave the users’ devices. Federated Learning (FL) (McMahan et al. 2023) with DP-FedAvg (McMahan et al. 2018) is commonly used to solve this problem. In Federated learning, clients locally train separate models on their data, that are then aggregated by a central server into a single, global model, which is then sent back to the clients to begin the next round of local training. Differential Privacy (DP) (Dwork 2006) can be applied to FL to allow the central server to obtain an effective global model without compromising user privacy. This paper focuses on Federated Learning as a platform for ensuring privacy, but each of our claims does not necessarily only apply to the FL case.

Robust federated learning. FL is vulnerable to attacks from malicious clients which construct local models that aim to introduce incorrect behaviour to the global model (Bagdasaryan et al. 2019). The most effective current methods for preventing these attacks remove local models that they determine lie outside benign model clusters. For example Krum discards models that lie far away from their closest neighbours in Euclidean space (Blanchard et al. 2017). The trimmed mean aggregation function instead discards the models with the n highest and lowest norms (Yin et al. 2021).

Alternatively, Sun et al. show that a weaker version of DP-FedAvg can improve robustness without directly rejecting specific clients. However, it has been shown that DP can have a disproportionate impact on clients that are further from the common model distribution (Bagdasaryan and Shmatikov 2019; Ganev, Oprisanu, and Cristofaro 2022; Farrand et al. 2020), so this method may indirectly have very similar effects to ‘anomaly detection’ methods such as Krum and trimmed mean.

Fair federated learning. In this paper, we define a model to be fair if it has an even accuracy distribution between different subpopulations of the dataset. For example, a phone manufacturer is likely to want face recognition accuracy to be equal between different ethnic groups. Fair aggregation functions typically attempt to increase the weights assigned to clients that hold less common types of data to increase the impact of these subsets of data. For example, in agnostic federated learning (Mohri, Sivek, and Suresh 2019), the aggregator returns the loss produced by the worst-case weighting of client loss values, while q-FFL (Li et al. 2020) weights clients by their loss, raised to the power of a parameter q .

Combining fair and robust FL methods. When attempting to construct a training procedure with both fair and robust methods, we encounter a contradiction: robust aggregation eliminates local models that lie outside the distribution of common models, while fair aggregation attempts to increase the effect of these models. After combining these techniques, we therefore expect that the benefits yielded by these methods may not entirely persist. In this paper, we

show that common robust aggregation methods can introduce unfairness during model training, and discuss how this affects our ability to prevent attacks on fairness. This is a fundamental problem with anomaly detection algorithms, which has so far not been directly addressed.

Related work. A similar idea has been previously presented by Wang et al., however their claims focus specifically on the implications for their proposed backdoor attack, while ours are broader. Additionally, there is previous work on constructing fair and robust training methods, however these methods either attempt to find a balance between fairness and robustness by using similar methods to those described above (Hu et al. 2022; Jin et al. 2023; Bowen et al. 2024), which we argue is not always an effective strategy, or solve the problem in non-standard setups, such as with personalisation (Li et al. 2021b), which are not applicable to the general case.

To our knowledge, no previous work has proposed attacks that target fairness in Federated Learning.

3 Attacks on Fairness in Federated Learning

In this section, we outline how fairness can be attacked in federated learning, and show that attacks on fairness are a practical threat in many realistic scenarios.

Threat model for attacks on fairness. We assume a very similar threat model to Bagdasaryan et al.: the attacker has access to the initial model sent to clients on the current round, G_t , and may control a minority subset of clients to submit arbitrary parameters for aggregation. **The attacker cannot see the models submitted by benign clients.** The server may view the submitted parameters, but does not know which are submitted by the attacker and which are from benign clients. Such a threat could reasonably exist in systems where clients are untrusted (i.e. the attacker joins as a new client), or can be compromised (e.g. by malware).

In an attack on fairness, the attacker aims to train the global model, G_{t+1} , to have high accuracy on some target dataset, $D_T \subseteq D$, and low accuracy on the other data, $D_N = D \setminus D_T$. For example, they may want high accuracy on only classes 0 and 1. The attacker can therefore obtain a set of target parameters, \mathbf{x} , that they wish to substitute into G_{t+1} , by fine-tuning G_t using only data in D_T .

The attacker then uses \mathbf{x} to compute a set of local, malicious parameters, \mathbf{c}_0 , such that, after aggregation with the other clients’ models, the resulting global model, G_{t+1} , is approximately equal to \mathbf{x} .

Model replacement attacks. If we directly submit \mathbf{x} to the aggregator (i.e. $\mathbf{c}_0 = \mathbf{x}$), our parameters are unlikely to have a significant effect after aggregation with a much larger volume of benign parameters. Bagdasaryan et al. propose a more powerful strategy, *model replacement*, that allows the attacker to substitute the arbitrary target parameters, \mathbf{x} , directly into the global model. Bagdasaryan et al. train \mathbf{x} to contain a *backdoor*, instead of introducing unfairness. Under the FedAvg (McMahan et al. 2023) aggregator, the attack by Bagdasaryan et al. is able to influence the global model to be $G_{t+1} = \mathbf{x} + \sum_{i=1}^{m-1} \frac{n_i}{n} (\mathbf{c}_i - G_t)$, where n_i

is the size of client i 's dataset, and $n = \sum_{i=0}^{m-1} n_i$.² Bagdasaryan et al. implicitly assume the model converges (and therefore $\mathbf{c}_i - G_t \approx 0$; see eqn. 3 in Bagdasaryan et al.), to obtain $G_{t+1} = \mathbf{x}$. However, unlike in the backdoor case, attacks on fairness prevent convergence.

The update prediction attack. The model replacement attack requires this convergence assumption because we do not know the parameters produced by other clients. For our attack on fairness, we solve this problem by letting the attacker *predict* the parameters submitted by other clients. If we assume that the difference between the mean client parameters, $\sum_{i=0}^m \frac{n_i}{n}(\mathbf{c}_i)$, and some other set of parameters, \mathbf{w} , that have been trained on data that is i.i.d. to the union of the clients' datasets is normally distributed with 0 mean, and variance a decreasing function of the amount of data seen during training, tending to 0 in the limit (We prove this is true below as **Theorem 1**), an attacker may be able to accurately predict the value $\sum_{i=0}^m \frac{n_i}{n}(\mathbf{c}_i)$. This forms the basis for the update prediction attack.

Instead of subtracting the global model, G_t , from our \mathbf{x} to yield \mathbf{c}_0 , as is the case in the model replacement attack, the attacker can now subtract their model predictions, \mathbf{w} , thus (approximately) eliminating *all* other terms in the FedAvg update rule. We set $\mathbf{c}_0 = \frac{n_0-n}{n_0}\mathbf{w} + \frac{n}{n_0}\mathbf{x}$ as the attacker's parameters, so the FedAvg update becomes

$$\begin{aligned} G_{t+1} &= G_t + \sum_{i=0}^{m-1} \frac{n_i}{n} (\mathbf{c}_i - G_t) \\ &= G_t + \frac{n_0}{n} \left(\frac{n_0-n}{n_0} \mathbf{w} + \frac{n}{n_0} \mathbf{x} - G_t \right) + \sum_{i=1}^m \frac{n_i}{n} (\mathbf{c}_i - G_t) \\ &= G_t + \frac{n_0-n}{n} \mathbf{w} + \mathbf{x} + \sum_{i=1}^m \frac{n_i}{n} (\mathbf{c}_i) - \sum_{i=0}^m \frac{n_i}{n} (G_t) \\ &= \mathbf{x} + \sum_{i=1}^m \frac{n_i}{n} (\mathbf{c}_i) - \frac{n-n_0}{n} \mathbf{w} \\ &\approx \mathbf{x} \end{aligned}$$

Here we directly get $G_{t+1} = \mathbf{x}$ without the convergence assumption, allowing us to use any set of parameters for \mathbf{x} (see fig. 2).

Following our threat model, **no knowledge of parameters submitted by other clients is required by this attack**. The attacker only requires an approximate estimate for the amount of data, $n - n_0$, contributed by other clients. Such an estimate need not be exact and could be iteratively increased each round until an effective value is found.

Attacks on fairness are well-founded. Our update prediction attack on fairness assumes that, when a federated learning model trains with a large amount of private data, the variance in benign client parameters introduced by unknown training data is low. We now prove this for strongly convex functions.

²We provide a full notation reference in appendix B




Figure 2: Visual representation of how the attack is constructed. Each arrow represents a single client's parameter vector. In practice, the angles between \mathbf{w} , and \mathbf{x} tend to be small, so the length of \mathbf{c}_0 is not as extreme as this diagram suggests.

We begin by considering the following general optimisation problem for client i :

$$\min_{\mathbf{c}_i \in \mathbb{R}^d} f(\mathbf{c}_i) = \min_{\mathbf{c}_i \in \mathbb{R}^d} \frac{1}{n_i} \sum_{j=0}^{n_i-1} f_j(\mathbf{c}_i) \quad (1)$$

where each $f_i \in \mathbb{R}^d \rightarrow \mathbb{R}$ is a continuously differentiable function. We want to use mini-batch gradient descent to solve this problem:

$$\mathbf{c}_i^{(k)} = \mathbf{c}_i^{(k-1)} - \frac{\alpha_k}{b} \sum_{j=0}^{b-1} \nabla f_{s_{k,j}}(\mathbf{c}_i^{(k-1)}) \quad (2a)$$

$$= \mathbf{c}_i^{(k-1)} - \alpha_k [\nabla f(\mathbf{c}_i^{(k-1)}) + \xi_k] \quad (2b)$$

where α_k is the learning rate, each $s_{k,j}$ is a uniformly random sample from $\{0, \dots, n_i - 1\}$, and ξ_k represents the noise introduced by sampling from the training distribution on round k . This problem description and the following assumptions follow that of Li, Xiao, and Yang, whose proof of SGD convergence to normally distributed models in the central case provides a basis for the following proofs. For simplicity, we set the learning rate to $\alpha_k = \alpha_1 k^{-1/2}$, which satisfies the assumptions required by Li, Xiao, and Yang.

We make the following assumptions in the below proofs.

(A1) Mean and covariance of ξ_k . $\forall \varepsilon > 0. \exists$ a symmetric, positive definite matrix, Σ , such that

$$\mathbb{E}[\xi_k | \mathcal{F}_{k-1}] = 0 = \lim_{n \rightarrow \infty} P(|\mathbb{E}[\xi_k \xi_k^T | \mathcal{F}_{k-1}] - \Sigma| \geq \varepsilon)$$

where $\mathcal{F}_k = \sigma(x_0, \xi_1, \xi_2, \dots, \xi_k)$ is the σ -algebra generated from the random initialisation and noise terms up to round k .

(A2) L -smoothness of f . $\exists L$ such that

$$\forall x, y \in \mathbb{R}^d. \|\nabla f(x) - \nabla f(y)\| \leq L\|x - y\|$$

(A3) μ -strong convexity of f . $\exists \mu$ such that

$$\forall x, y \in \mathbb{R}^d. f(x) \geq f(y) + \nabla f(y)^T (x - y) + \frac{\mu}{2} \|x - y\|^2$$

(A4) Further smoothness condition for f . $\exists p_0, r_0, K_d > 0$ such that, for any $\|x - x^*\| < r_0$,

$$\|\nabla f(x) - \nabla^2 f(x^*)(x - x^*)\| \leq K_d \|x - x^*\|^{1-p_0}$$

(A5) Dataset size heterogeneity. If client i has a dataset of size n_i , modelled as a random variable, and $n = \sum_{c=0}^{m-1} n_i$ for m clients, then,

$$\exists h \geq 1 \in \mathbb{R}. \forall i \in [0, 1, \dots, m-1]. mn_i \leq nh$$

With this definition, if $h = 1$, all client datasets must have the same amount of data, while as $h \rightarrow \infty$, the client data distribution constraints disappear.

Lemma 1 (variance for a single client). Under assumptions (A1)-(A4), if $\frac{1}{\alpha_1} < 2\mu$, there exists some matrix, W^* , such that

$$k^{1/4}(\mathbf{c}_i^{(k)} - \mathbf{c}_i^*) \Rightarrow^k N(0, \alpha_1 W^*) \quad (3)$$

where \Rightarrow^k denotes convergence in probability, \mathbf{c}_i^* is the unique minimum of f , $k \in \mathbb{N}$ is large, and $W_{k,i,j}^* \in O\left(\frac{1}{b^2}\right)$.

Proof. This extends the result from Li, Xiao, and Yang. Under the above assumptions, Li, Xiao, and Yang show that for $b = 1$ we have eq. (3) for some matrix W^* , where $AW^* + W^*A^T - d_0W^* = \Sigma$ and A is independent of all ξ_i .

Now consider the variance of ξ_k as b increases. We assume that the summed gradients are independent and have finite first 2 moments. Thus, for large b , by the classical CLT, the gradient estimate, $\nabla \hat{f}(\mathbf{c}_i^{(k-1)})$, is unbiased and normally distributed:

$$\begin{aligned} \nabla \hat{f}(\mathbf{c}_i^{(k-1)}) &= \frac{1}{b} \sum_{j=0}^{b-1} \nabla f_{s_{k,j}}(\mathbf{c}_i^{(k-1)}) \\ &\sim N\left(\mathbb{E}_{f_j}[\nabla f_j(\mathbf{c}_i^{(k-1)})], \frac{\mathbb{V}_{f_j}[f_j(\mathbf{c}_i^{(k-1)})]}{b}\right) \end{aligned}$$

This yields a noise term, $\xi_k \sim N\left(0, \mathbb{V}[f_j(\mathbf{c}_i^{(k-1)})]/b\right)$, with variance inversely proportional to batch size.

The maximum element in the covariance matrix for $\text{vec}(\xi_k \xi_k^T)$ (where vec is a function that flattens a matrix into a vector) must be the variance of $(\xi_k)_i^2$ for some i . Since $(\xi_k)_i$ is normally distributed with variance $\frac{c}{b}$ for some constant c , we know that each element of this covariance matrix must be bounded by $\mathbb{V}_{\xi_k}[(\xi_k)_i^2] = \frac{2c^2}{b^2}$.

We have established that $AW^* + W^*A^T - d_0W^* = \Sigma$ and that the elements of the covariance matrix for $\xi_k \xi_k^T$ (and therefore also those of Σ) are in $O\left(\frac{1}{b^2}\right)$, so the elements of W^* must also be in $O\left(\frac{1}{b^2}\right)$. \square

Now consider the FedAvg aggregation function (McMahan et al. 2023) to compute the global model, G , from the model $\mathbf{c}_i^{(u)}$ produced by each client i after u batches using the above SGD setup for the current training round:

$$G = \frac{1}{n} \sum_{i=0}^{m-1} n_i \mathbf{c}_i^u \quad (4)$$

where $n = \sum_{i=0}^m n_i$.

Theorem 1 (Variance of FedAvg). Under assumptions (A1)-(A5), if $\frac{1}{\alpha_1} < 2\mu$ for each client, the global model, G , must be normally distributed with covariance matrix M_g such that $M_{g,p,q} \in O\left(1/\sqrt[4]{enm^3b^7}\right)$ for a large epoch number, e , and batch size, b .

Proof. From **Lemma 1**, we know that each $\mathbf{c}_i^{(e-1)}$ is an independent, normally distributed random variable with covariance matrix M_i , where $M_{i,p,q} \in O\left(1/\sqrt[4]{ub^8}\right) = O\left(1/\sqrt[4]{en_ib^7}\right)$, for large e and b . By applying the FedAvg procedure, we get

$$G \sim N\left(\sum_{i=0}^{m-1} \frac{n_i}{n} \mathbf{c}_i^*, \sum_{i=0}^{m-1} \frac{n_i^2}{n^2} M_i\right) \quad (5)$$

Since, by (A5), $\max_i \frac{n_i}{n} \in O\left(\frac{1}{m}\right)$ for all clients, i , the covariance matrix $M_g = \sum_{i=0}^{m-1} \frac{n_i^2}{n^2} M_i$ must have $M_{g,p,q} \in O\left(1/\sqrt[4]{en_im^4b^7}\right) = O\left(1/\sqrt[4]{enm^3b^7}\right)$. \square

Therefore, by Chebyshev's inequality, **the probability of our model prediction error being greater than γ is bounded by $O\left(1/\sqrt[4]{enm^3b^7\gamma^8}\right)$.**

The above proof can similarly be adapted for SGD with momentum, using the relevant results from Li, Xiao, and Yang. This proof clearly does not extend to all non-convex models. However, similarly to model convergence in general, it is reasonable to assume that for a sufficiently smooth loss function and large enough batch size, because the attacker knows the model's initial parameters, the non-convex case is locally similar to the strongly convex case above.

The update prediction attack also assumes that \mathbf{w} is an unbiased estimator of $\sum_{i=0}^m \frac{n_i}{n} (\mathbf{c}_i)$. This is false when there is heterogeneity between clients (FedAvg introduces some unfairness itself). The attacker could construct an unbiased estimator by directly locally simulating the entire FL training process, however we find that predicting \mathbf{w} in a centralised manner is effective in practice.

Experimental results. We now test the fairness attack for each of the three datasets described in table 2 (Becker and Kohavi 1996; Krizhevsky 2009; Pushshift 2017). These datasets were selected to cover a range of tasks and to provide clear comparison with previous work (Bagdasaryan et al. 2019; Bhagoji et al. 2019; Wang et al. 2020; Nguyen et al. 2023; McMahan et al. 2023). Here, we substitute a single client's model with our malicious set of parameters, using the original client's model as the prediction \mathbf{w} (i.e. we predict \mathbf{w} using the same amount of data that would be held on a single client). For simplicity, we include the attack in every round. We additionally test its performance against the Krum, trimmed mean, and weak differential privacy defences. We select hyperparameters by performing a grid search over all reasonable combinations at multiple levels of granularity and present the median result across three trials in table 1. All experiments were performed on 2 NVIDIA RTX 2080 GPUs.

We record the change in fairness after the attack is introduced for each dataset-defence combination. The attack is effective at introducing unfairness into all three tasks. In practice, it may be preferable to perform a more subtle version of this attack. Although the size of the dataset needed to train \mathbf{w} depends on the task, the attack remains effective even with the small local datasets available in the Reddit task. Table 2 shows that the Krum and trimmed mean defences are effective at preventing the attack on fairness, however we find that the weak DP defence was not successful under any hyperparameter configuration.

Dataset	Attack	No defence		Krum		Trimmed mean		Weak DP	
		Acc.	Δ fairness	Acc.	Δ fairness	Acc.	Δ fairness	Acc.	Δ fairness
Census	None	84.81	0.00	84.82	0.00	84.83	0.00	75.66	0.00
	Fairness	81.57	53.44	84.64	-1.72	84.76	1.82	76.38	2.23
CIFAR-10	None	92.70	0.00	91.59	0.00	91.98	0.00	93.60	0.00
	Fairness	17.90	85.24	63.78	6.26	63.11	6.28	17.62	85.29
Reddit	None	18.08	0.00	17.82	0.00	18.08	0.00	08.55	0.00
	Fairness	4.52	118.94	17.97	-0.16	17.90	0.19	04.52	108.39

Table 1: Accuracy (%) achieved by different robust aggregation schemes for each dataset. The attack on fairness attempts to increase the accuracy on one subset of data while reducing the accuracy on another, so ‘ Δ fairness’ indicates the increase in accuracy disparity between these two sets (lower is better).

Dataset	#Records	Model	#Clients/Per round
UCI Census	49k	3-layer FC	10 / 10
CIFAR-10	60k	ResNet-50	10 / 10
Reddit	2.3M	LSTM	10,000 / 100

Table 2: We train clients for 10, 2, and 5 epochs on i.i.d. data, for a total of 40, 120, and 100 rounds for the Census, CIFAR, and Reddit datasets respectively. We use the same augmentation scheme as Zagoruyko and Komodakis for the CIFAR-10 dataset and the *albert-base-v2* tokeniser (Lan et al. 2020) for the Reddit task. In the Census task, the attack reduces the accuracy of entries labelled as female, in the CIFAR-10 task, it increases accuracy on classes 0 and 1, and in the Reddit task, it reduces accuracy following the word ‘I’.

Previous work has questioned the effectiveness of the weak differential privacy defence (Wang et al. 2020) in preventing adversarial attacks, although it is difficult to justify specifically why the defence is ineffective against this attack-task combination. However, recent works have also shown that weak differential privacy can introduce unfairness into a model (Bagdasaryan and Shmatikov 2019; Ganey, Oprisanu, and Cristofaro 2022; Farrand et al. 2020).

Attacks against momentum-based aggregation functions. Momentum-based aggregation functions (Reddi et al. 2021) are common in practical applications, which could lead to the attack on fairness becoming ineffective. We therefore repeat our experiments for the Census task in table 1 for three momentum-based aggregators (table 3), finding that attacks on fairness remain effective against these aggregators without modification.

Aggregator	Attack	Overall	Δ fairness
FedAdaGrad	Baseline	84.58	-2.84
	Fairness	75.79	36.47
FedYogi	Baseline	78.72	43.76
	Fairness	80.78	51.87
FedAdam	Baseline	84.50	1.64
	Fairness	80.14	45.91

Table 3: Attacks on different aggregators for the Census task. All hyperparameters were identical to table 1, which accounts for the drop in initial accuracy and fairness. Δ fairness is calculated with respect to the baseline in table 1.

Attacks on fairness under high data heterogeneity. We have shown that attacks on fairness function without significant data heterogeneity between clients. However, in most real tasks, this is not the case (Yang et al. 2018; Hard et al. 2019; Huba et al. 2022). Furthermore, by tracing the effect of increasing h in (A5), we expect heterogeneity to reduce the attack’s strength.

To show that attacks on fairness can be a threat in settings where there is high heterogeneity, we repeat the baseline experiment from table 1, with the CIFAR-10 dataset distributed between clients using a log-normal label distribution across the clients parameterised by $\mu = 0$ and $\sigma \in \{0, 1, 2\}$. Figure 3 shows that heterogeneity reduces the attack’s effectiveness. However, even at high σ values, it remains effective. Here, we do not include any defence, although we should expect that increased heterogeneity would make detection more difficult (Ozdayi and Kantarcioglu 2021).




Figure 3: Accuracy per training round for the attack on fairness under different heterogeneity values (σ) for the CIFAR-10 task. The red (upper) line represents the accuracy on data the attack seeks to increase, while the green (lower) line represents the accuracy on other data. As heterogeneity increases, the attack’s strength decreases. Without the fairness attack, the accuracy on both datasets is approximately equal.

Why an attacker might target fairness. While the existence of attacks on fairness is sufficient for many of the claims in this paper, we might question whether clients should reasonably be expected to employ these attacks. Here we provide one example where a client may wish to attack the fairness of a model.

Consider a model trained to predict interactions between pairs of molecules (Mittone et al. 2023). The model is trained using data on drugs produced by a group of participating companies, with federated learning employed to ensure that these companies’ proprietary training data is not leaked. One company may wish to employ an attack on fairness on the model to reduce the accuracy of interactions with drugs produced by competitors, in order to gain an advantage. In this scenario, we are faced with the dilemma: either we risk vulnerability to this attack on fairness, or we risk discarding data on rare drug interactions that are misidentified as attacks.

4 Robust Aggregation Introduces Unfairness

In the previous section we have shown that attacks on fairness are a realistic threat during federated learning training. This motivates the need for a training process that is robust to these attacks, without introducing unfairness itself. In this section we seek to answer the question **can we prevent adversarial attacks on federated learning without introducing unfairness?**

Fairness impact of Krum and trimmed mean. Both the Krum and trimmed mean robust aggregation methods remove client models that lie far from a mode of the distribution. This could lead to unfairness because clients holding certain types of uncommon data are more likely to produce models that lie far from clusters of common models than other clients (see fig. 4). Figure 5 shows that some clients hold meaningfully less common datasets even when there is low data heterogeneity between clients (some clients are consistently ranked more trustworthy than others when data is randomly distributed between them). Furthermore, removal of a small number of these clients can have a disproportionately high impact due to the uncommon nature of the data they hold.

We now show that for both defences, this can lead to the reduction, and, in this extreme case, elimination, of critical functionality from a model when this functionality is only present in a minority of clients.




Figure 4: 2D projection of models produced by clients on the MNIST dataset (LeCun, Cortes, and Burges 1998). Points coloured green (inverted triangles) are produced by clients that only hold data with classes 0 and 1, while points coloured red (upright triangles) are produced by clients running a backdoor attack. Although a low-dimensional representation exaggerates the problem, we can see that there is little difference between malicious models and models trained on specific subpopulations of the dataset from a clustering perspective.

We construct the dataset shown in fig. 6, in which five clients (group A) do not have any data where the input begins with a 0, and one client (group B) does not have data where the input ends with a 0. This construction leads to the client from group B producing different models compared to group A when training on a simple, fully-connected model.

As shown in table 4, both defence methods incorrectly determine that the group B client is malicious across multiple tests, leading to the uncommon functionality that is unique to this client ($[0, 1] \rightarrow 1$) being lost in the aggregated model. As this data is lost, fair aggregation methods that reweight uncommon models would be ineffective in this scenario. Although we eliminate more models (1) than there are attackers (0) on each round, this is realistic because we will not know how many attackers there are in practice.



Figure 5: Client trustworthiness ranking on each round according to the Krum defence for the CIFAR-10 task. The attacker (red) is consistently ranked least trustworthy, however, we also observe some clients (e.g. the lower line coloured blue) are consistently lower ranked than others (e.g. the upper line coloured green).

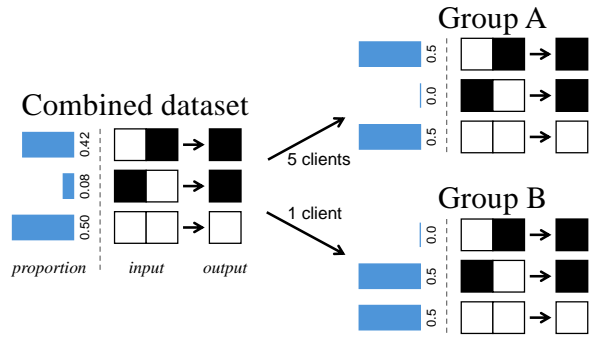


Figure 6: This dataset represents the OR function, where black squares indicate the value 1, and white the value 0. We split it so that the input $[0, 1]$ only occurs in 5 out of 6 clients, while the input $[1, 0]$ only occurs in the remaining client.

Furthermore, the aggregator does not know how each local model is obtained (due to our privacy constraints) so there cannot exist *any* unsupervised anomaly-detection-based defence that would be able to identify group B as benign, while still rejecting all attacks. This is true because, if such a defence existed, it would continue to accept models from group B when the task is redefined as returning the value of the second input (i.e. $[0, 1] \rightarrow 1$; $[0, 0] \rightarrow 0$; \dots). In this scenario, group B introduces behaviour that directly opposes the training goal, so it should reasonably be classed as malicious.

More generally, the local training function that produces models from local data by SGD is non-injective and thus has no left inverse, so overlap between the tail of the benign distribution and the set of malicious models is possible (see fig. 7). Shumailov et al. show that malicious parameters can be learnt by manipulating the order of clean data, which implies that this overlap exists in realistic scenarios. **Therefore, even without differential privacy constraints, for some datasets it is impossible to detect all attacks without misidentifying some benign models as malicious.**

In our testing shown in table 1, we also find that accuracy disparity between common and uncommon data in the Census task³

³In our other two tasks, there is no clear dataset split that yields such uneven subpopulations

Defence	Combined	Group A	Group B
No defence	100	100	100
Trimmed mean	92	100	50
Krum	92	100	50

Table 4: Accuracy (%) for each dataset under different defences.




Figure 7: There may be overlap between the distribution of models produced by legitimate clients (blue, left) and the set of models produced by malicious processes (red, right). Methods based on anomaly detection must reject the tail of the benign distribution (e.g. by accepting only models to the left of boundary A, leading to unfairness due to the omission of uncommon data) or accept some malicious models (e.g. by accepting only models to the left of boundary B).

increased under all three defences, with the disparity growing by a median of 1.95, 0.53, and 51.72 for the trimmed mean, Krum, and weak DP defences respectively. Similar results to these have previously been observed in realistic setups, often showing a more significant reduction in fairness compared to these tasks (for example, Wang et al. 2020), although analysis of this issue has never been extended to the general problem that we investigate here. Thus, even without the use of techniques that attempt evade detection by robust aggregators (Bagdasaryan et al. 2019), which can serve to increase the difficulty of separating uncommon from malicious models, unfairness is introduced due to this overlap problem.

While robust aggregation methods that attempt to retain fairness have been presented, they are forced to make significant compromises compared to the methods studied here. For example, FedMGDA+ (Hu et al. 2022) employs a gradient clipping strategy, which is clearly weaker than the weak differential privacy defence described above.

Unfair-update detection: testing a new defence for fairness attacks. While Wang et al. have shown that verifying that a model does not contain any backdoors is computationally intractable, it is relatively simple to verify that a model is fair across a set of predetermined attributes with high confidence. This suggests a simpler defence for attacks that only attempt to introduce unfairness may be to measure the fairness impact of each client’s model and assume clients that significantly reduce fairness are malicious. This defence is described more specifically by algorithm 1 (appendix A).

When repeating the experiments for the CIFAR-10 task with this defence, the change in accuracy disparity (Δ fairness) is only 0.34 under the fairness attack. Additionally, the unfair-update detection algorithm initially appears to solve the problem of overlapping malicious and benign model distributions by accepting models based on their impact on the global model rather than based on how we believe they have been trained. However, a client which submits a model trained on new data that may be necessary to achieve a more fair *final* model (after convergence) is likely to have reduced accuracy/fairness in the short term, leading to its rejection in some sce-

narios. The problem arises from the attack’s greedy setup: we select the clients that produce the most fair model *on the next round*, not those that result in a more fair *final* model, when it may be necessary to temporarily reduce fairness in order to achieve a higher final value.

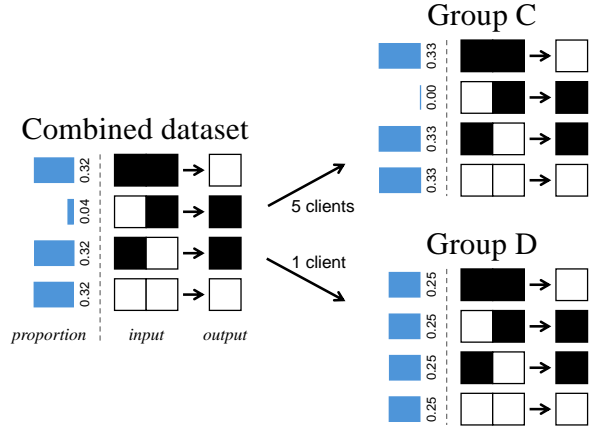


Figure 8: This dataset represents the XOR function. We split it so that the input $[0, 1]$ only occurs in 1 out of 6 clients.

To demonstrate that this defence does not represent a solution to the problem faced by the Krum and trimmed mean defences, we construct the dataset shown in fig. 8 to show that there are some tasks for which it will, counterintuitively, *increase* unfairness. To maximise fairness, we must include group D, but introducing group D for a single round reduces fairness by producing weights that reduce the accuracy of some instances of group C data, leading to the group D client’s rejection by the defence. When training on this dataset, a simple, two-layer, fully-connected model shows significantly lower accuracy for group D with the defence compared to without (table 5).

Defence	Combined	Group C	Group D
No defence	100	100	100
Unfair-update det.	96	100	75

Table 5: Accuracy (%) for each dataset under different defences.

5 Conclusion

We have shown that common defence methods can introduce unfairness and that attacks on fairness are a real threat to the federated learning training process. Furthermore, to our knowledge, there does not exist any defence that can ensure the robustness of the global model without using a method based on those analysed in this paper. **Thus, assuming such a defence does not currently exist, in the presence of untrusted clients, we cannot be confident that training on private data will result in a fair model.**

Even if robust aggregation may introduce a negligible amount of unfairness for *some* datasets, it is difficult to predict which datasets will present this problem and how to manage the tradeoff between fairness and robustness. This presents a large risk, particularly for systems that are expensive to retrain. Future work would therefore benefit from approaching the problem of robustness in FL from a new, fairness-aware perspective.

A Unfair-update Detection Pseudocode

Algorithm 1: Unfair-update detection

Input: number of malicious clients, a ; vector of datasets to compute fairness across, \mathbf{d}

Output: global model at round t

```

Initialise global model  $G_0$ 
for each training round,  $t$ , in  $[1, T]$  do
  for each client,  $i$ , in  $[1, m]$  do
     $r_i \leftarrow \text{get\_client\_update}(G_{t-1})$  {server sends client  $i$ 
       $G_{t-1}$  and gets response  $r_i$ }
  end for
  for each combination of  $m - a$  clients,  $c$  do
     $G'_t \leftarrow \text{aggregate}(\{r_i | i \in c\})$ 
     $f \leftarrow \text{compute\_fairness}(G'_t, \mathbf{d})$ 
    if  $f$  is the most fair so far then
       $G_t \leftarrow G'_t$ 
    end if
  end for
end for

```

B Notation Reference

G_t	The parameter vector of the global model at time t (uppercase used for consistency with previous work)
\mathbf{c}_i	Client i 's parameter vector at the current round
m	The total number of clients in the current round
n_i	The number of training samples owned by client i
n	The total amount of data owned by all clients
h	The heterogeneity factor of the participating clients
\mathbf{x}	The parameters an attacker aims to replace G_{t+1} with
D_T	A dataset containing only data the fairness attack aims to increase the accuracy of
D_N	A dataset containing only data the fairness attack aims to decrease the accuracy of
\mathbf{w}	The attacker's prediction of the aggregated model without the malicious parameters
u	The number of batches each local model is trained for
e	The number of epochs each local model is trained for
k	The current epoch
α_k	The local learning rate at epoch k
$\mathbf{c}_i^{(k)}$	The estimate of client i 's parameters at epoch k
b	The mini-batch size
$s_{k,j}$	The j th uniformly random data sample at epoch k
ξ_k	The noise from sampling from \mathcal{D}_i on round k
μ	The loss function's convexity term
L	The loss function's smoothness term
\mathbf{c}_i^*	The unique minimum of the loss function for client i
$M_{i,j}$	The value at row i and column j of matrix M .
$\nabla f(x)$	The gradient of f at point x
$\nabla \hat{f}(x)$	A normally dist., unbiased estimate of $\nabla f(x)$
$P(f(X))$	The probability f is true for R.V. X
\mathbb{E}_X	The expectation of $f(X)$ over R.V. X
\mathbb{V}_X	The variance of $f(X)$ over R.V. X
$N(\mu, \sigma^2)$	A normal dist. with mean μ and variance σ^2
$X_k \Rightarrow^k X$	X_k converges in probability to X as $k \rightarrow \infty$

Ethical Statement

This work introduces a new type of attack that could be used to maliciously affect real systems. We provide a sample implementation of this attack, because we believe the benefits of improving the accessibility of tools to test the robustness of existing and future systems outweigh the disadvantages of making these attacks more accessible to people who may use them maliciously.

References

- Arachchige, P. C. M.; Bertók, P.; Khalil, I.; Liu, D.; Çamtepe, S. A.; and Atiquzzaman, M. 2020. A Trustworthy Privacy Preserving Framework for Machine Learning in Industrial IoT Systems. *IEEE Transactions on Industrial Informatics*, 16: 6092–6102.
- Bagdasaryan, E.; and Shmatikov, V. 2019. Differential Privacy Has Disparate Impact on Model Accuracy. arXiv:1905.12101.
- Bagdasaryan, E.; Veit, A.; Hua, Y.; Estrin, D.; and Shmatikov, V. 2019. How To Backdoor Federated Learning. arXiv:1807.00459.
- Becker, B.; and Kohavi, R. 1996. Adult. UCI Machine Learning Repository.
- Bhagoji, A. N.; Chakraborty, S.; Mittal, P.; and Calo, S. 2019. Analyzing Federated Learning through an Adversarial Lens. arXiv:1811.12470.
- Blanchard, P.; El Mhamdi, E. M.; Guerraoui, R.; and Stainer, J. 2017. Machine Learning with Adversaries: Byzantine Tolerant Gradient Descent. In Guyon, I.; Luxburg, U. V.; Bengio, S.; Wallach, H.; Fergus, R.; Vishwanathan, S.; and Garnett, R., eds., *Advances in Neural Information Processing Systems*, volume 30. Curran Associates, Inc.
- Bowen, D.; Haiquan, W.; Yuxuan, L.; Zhao, J.; Ma, Y.; and Runhe, H. 2024. Fair and Robust Federated Learning via Decentralized and Adaptive Aggregation based on Blockchain. *ACM Trans. Sen. Netw.* Just Accepted.
- Cao, X.; Fang, M.; Liu, J.; and Gong, N. Z. 2020. FLTrust: Byzantine-robust Federated Learning via Trust Bootstrapping. *ArXiv*, abs/2012.13995.
- Cho, Y. J.; Wang, J.; and Joshi, G. 2020. Client Selection in Federated Learning: Convergence Analysis and Power-of-Choice Selection Strategies. arXiv:2010.01243.
- Dwork, C. 2006. Differential Privacy. In Bugliesi, M.; Preneel, B.; Sassone, V.; and Wegener, I., eds., *Automata, Languages and Programming*, 1–12. Berlin, Heidelberg: Springer Berlin Heidelberg. ISBN 978-3-540-35908-1.
- Farrand, T.; Mireshghallah, F.; Singh, S.; and Trask, A. 2020. Neither Private Nor Fair: Impact of Data Imbalance on Utility and Fairness in Differential Privacy. arXiv:2009.06389.
- Ganev, G.; Oprisanu, B.; and Cristofaro, E. D. 2022. Robin Hood and Matthew Effects: Differential Privacy Has Disparate Impact on Synthetic Data. arXiv:2109.11429.
- Hard, A.; Rao, K.; Mathews, R.; Ramaswamy, S.; Beaufays, F.; Augenstein, S.; Eichner, H.; Kiddon, C.; and Ramage, D. 2019. Federated Learning for Mobile Keyboard Prediction. arXiv:1811.03604.
- Hu, Z.; Shaloudegi, K.; Zhang, G.; and Yu, Y. 2022. Federated Learning Meets Multi-Objective Optimization. *IEEE Transactions on Network Science and Engineering*, 9(4): 2039–2051.
- Huba, D.; Nguyen, J.; Malik, K.; Zhu, R.; Rabbat, M.; Yousefpour, A.; Wu, C.-J.; Zhan, H.; Ustinov, P.; Srinivas, H.; Wang, K.; Shoumikhin, A.; Min, J.; and Malek, M. 2022. Papaya: Practical, Private, and Scalable Federated Learning. arXiv:2111.04877.

- Jin, S.; Li, Y.; Chen, X.; Li, R.; and Shen, Z. 2023. Blockchain-Based Fairness-Enhanced Federated Learning Scheme Against Data Poisoning Attack. In *Smart Computing and Communication: 7th International Conference, SmartCom 2022, New York City, NY, USA, November 18–20, 2022, Proceedings*, 329–339. Berlin, Heidelberg: Springer-Verlag. ISBN 978-3-031-28123-5.
- Krizhevsky, A. 2009. Learning Multiple Layers of Features from Tiny Images.
- Lan, Z.; Chen, M.; Goodman, S.; Gimpel, K.; Sharma, P.; and Soricut, R. 2020. ALBERT: A Lite BERT for Self-supervised Learning of Language Representations. arXiv:1909.11942.
- LeCun, Y.; Cortes, C.; and Burges, C. J. 1998. The mnist database of handwritten digits.
- Li, T.; Beirami, A.; Sanjabi, M.; and Smith, V. 2021a. Tilted Empirical Risk Minimization. arXiv:2007.01162.
- Li, T.; Hu, S.; Beirami, A.; and Smith, V. 2021b. Ditto: Fair and Robust Federated Learning Through Personalization. arXiv:2012.04221.
- Li, T.; Sanjabi, M.; Beirami, A.; and Smith, V. 2020. Fair Resource Allocation in Federated Learning. arXiv:1905.10497.
- Li, T.; Xiao, T.; and Yang, G. 2023. Revisiting the central limit theorems for the SGD-type methods. arXiv:2207.11755.
- McMahan, H. B.; Moore, E.; Ramage, D.; Hampson, S.; and y Arcas, B. A. 2023. Communication-Efficient Learning of Deep Networks from Decentralized Data. arXiv:1602.05629.
- McMahan, H. B.; Ramage, D.; Talwar, K.; and Zhang, L. 2018. Learning Differentially Private Recurrent Language Models. arXiv:1710.06963.
- Mittone, G.; Svoboda, F.; Aldinucci, M.; Lane, N.; and Lió, P. 2023. A Federated Learning Benchmark for Drug-Target Interaction. In *Companion Proceedings of the ACM Web Conference 2023, WWW '23 Companion*, 1177–1181. New York, NY, USA: Association for Computing Machinery. ISBN 9781450394192.
- Mohri, M.; Sivek, G.; and Suresh, A. T. 2019. Agnostic Federated Learning. arXiv:1902.00146.
- Nguyen, T. D.; Rieger, P.; Chen, H.; Yalame, H.; Möllering, H.; Fereidooni, H.; Marchal, S.; Miettinen, M.; Mirhoseini, A.; Zeitouni, S.; Koushanfar, F.; Sadeghi, A.-R.; and Schneider, T. 2023. FLAME: Taming Backdoors in Federated Learning (Extended Version 1). arXiv:2101.02281.
- Nicola, R.; Jonny, H.; Wenqi, L.; Fausto, M.; Holger, R.; Shadi, A.; Spyridon, B.; Mathieu, G.; Bennett, L.; Klaus, M.; and S. Micah, O. S.; Ronald, S.; Andrew, T.; Daguang, X.; Maximilian, B.; and Jorge, C. 2020. The future of digital health with federated learning.
- Ozdayi, M. S.; and Kantarcioglu, M. 2021. The Impact of Data Distribution on Fairness and Robustness in Federated Learning. In *2021 Third IEEE International Conference on Trust, Privacy and Security in Intelligent Systems and Applications (TPS-ISA)*, 191–196.
- Pushshift. 2017. Reddit Comment Data. <https://files.pushshift.io/reddit/comments/>.
- Reddi, S.; Charles, Z.; Zaheer, M.; Garrett, Z.; Rush, K.; Konečný, J.; Kumar, S.; and McMahan, H. B. 2021. Adaptive Federated Optimization. arXiv:2003.00295.
- Shumailov, I.; Shumaylov, Z.; Kazhdan, D.; Zhao, Y.; Papernot, N.; Erdogdu, M. A.; and Anderson, R. 2021. Manipulating SGD with Data Ordering Attacks. arXiv:2104.09667.
- Soltan, A. A. S.; Thakur, A.; Yang, J.; Chauhan, A.; D’Cruz, L. G.; Dickson, P.; Soltan, M. A.; Thickett, D. R.; Eyre, D. W.; Zhu, T.; and Clifton, D. A. 2024. A scalable federated learning solution for secondary care using low-cost microcomputing: privacy-preserving development and evaluation of a COVID-19 screening test in UK hospitals.
- Sun, Z.; Kairouz, P.; Suresh, A. T.; and McMahan, H. B. 2019. Can You Really Backdoor Federated Learning? arXiv:1911.07963.
- Wang, H.; Sreenivasan, K.; Rajput, S.; Vishwakarma, H.; Agarwal, S.; yong Sohn, J.; Lee, K.; and Papailiopoulos, D. 2020. Attack of the Tails: Yes, You Really Can Backdoor Federated Learning. arXiv:2007.05084.
- Yang, T.; Andrew, G.; Eichner, H.; Sun, H.; Li, W.; Kong, N.; Ramage, D.; and Beaufays, F. 2018. Applied Federated Learning: Improving Google Keyboard Query Suggestions. arXiv:1812.02903.
- Yin, D.; Chen, Y.; Ramchandran, K.; and Bartlett, P. 2021. Byzantine-Robust Distributed Learning: Towards Optimal Statistical Rates. arXiv:1803.01498.
- Zagoruyko, S.; and Komodakis, N. 2017. Wide Residual Networks. arXiv:1605.07146.